

Practical Machine Learning Project - Quantified Self Movement Data Analysis Report

Venkatesh Attinti

4/22/2020

Practical Machine Learning Course Project Report

Background

Using devices such as Jawbone Up, Nike FuelBand, and Fitbit it is now possible to collect a large amount of data about personal activity relatively inexpensively. These type of devices are part of the quantified self movement - a group of enthusiasts who take measurements about themselves regularly to improve their health, to find patterns in their behavior, or because they are tech geeks. One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it. In this project, our goal will be to use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants. They were asked to perform barbell lifts correctly and incorrectly in 5 different ways. More information is available from the website here: <http://groupware.les.inf.puc-rio.br/har> (see the section on the Weight Lifting Exercise Dataset).

Data Sources

The training data for this project is available here:

<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>

The test data is available here:

<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>

The data for this project comes from this original source: <http://groupware.les.inf.puc-rio.br/har>. If you use the document you create for this class for any purpose please cite them as they have been very generous in allowing their data to be used for this kind of assignment.

Intended Results

The goal of this project is to predict the manner in which they did the exercise. This is the “classe” variable in the training set. You may use any of the other variables to predict with. You should create a report describing how you built your model, how you used cross validation, what you think the expected out of sample error is, and why you made the choices you did. You will also use your prediction model to predict 20 different test cases.

1. Your submission should consist of a link to a Github repo with your R markdown and compiled HTML file describing your analysis. Please constrain the text of the writeup to < 2000 words and the number of figures to be less than 5. It will make it easier for the graders if you submit a repo with a gh-pages branch so the HTML page can be viewed online (and you always want to make it easy on graders :-).
2. You should also apply your machine learning algorithm to the 20 test cases available in the test data above. Please submit your predictions in appropriate format to the programming assignment for automated grading. See the programming assignment for additional details.

Reproducibility

Below are packages used and the seed value, you need to use the same version of package to reproduce the results

: To install, for instance, the `rattle` package in R, run this command: `install.packages("rattle")`. The following Libraries were used for this project, which you should install and load them in your working environment.

```
library(rattle)
```

```
## Rattle: A free graphical interface for data science with R.  
## Version 5.3.0 Copyright (c) 2006-2018 Togaware Pty Ltd.  
## Type 'rattle()' to shake, rattle, and roll your data.
```

```
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
library(rpart)  
library(rpart.plot)  
library(corrplot)
```

```
## corrplot 0.84 loaded
```

```
library(randomForest)
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##  
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':  
##  
##     margin
```

```
## The following object is masked from 'package:rattle':  
##  
##     importance
```

```
library(RColorBrewer)
```

Setting the seed :

```
set.seed(56789)
```

Getting Data

Set your current working directory.

```
setwd("H:/RProjects/Practical Machine Learning")
```

The following code fragment downloads the dataset to the `data` folder in the current working directory.

```
trainUrl <- "https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv"
testUrl <- "https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv"
if (!file.exists("./data")) {
  dir.create("./data")
}

trainFile <- "./data/pml-training.csv"
testFile <- "./data/pml-testing.csv"

if (!file.exists(trainFile)) {
  download.file(trainUrl, destfile = trainFile)
}
if (!file.exists(testFile)) {
  download.file(testUrl, destfile = testFile)
}
```

Reading Data

After downloading the data from the data source, we can read the two csv files into two data frames.

```
trainRaw <- read.csv(trainFile)
testRaw <- read.csv(testFile)
dim(trainRaw)
```

```
## [1] 19622 160
```

```
dim(testRaw)
```

```
## [1] 20 160
```

The training data set contains `{r} dim(trainRaw)[1]` observations and `{r} dim(trainRaw)[2]` variables, while the testing data set contains `{r} dim(testRaw)[1]` observations and `{r} dim(testRaw)[2]` variables. The `classe` variable in the training set is the outcome to predict.

Cleaning Data

In this step, we will clean the dataset and get rid of observations with missing values as well as some meaningless variables.

1. Removing the Near Zero Variance Variables.

```
NZV <- nearZeroVar(trainRaw, saveMetrics = TRUE)
head(NZV, 20)
```

##		freqRatio	percentUnique	zeroVar	nzv
## X		1.000000	100.00000000	FALSE	FALSE
## user_name		1.100679	0.03057792	FALSE	FALSE
## raw_timestamp_part_1		1.000000	4.26562022	FALSE	FALSE
## raw_timestamp_part_2		1.000000	85.53154622	FALSE	FALSE
## cvtd_timestamp		1.000668	0.10192641	FALSE	FALSE
## new_window		47.330049	0.01019264	FALSE	TRUE
## num_window		1.000000	4.37264295	FALSE	FALSE
## roll_belt		1.101904	6.77810621	FALSE	FALSE
## pitch_belt		1.036082	9.37722964	FALSE	FALSE
## yaw_belt		1.058480	9.97349913	FALSE	FALSE
## total_accel_belt		1.063160	0.14779329	FALSE	FALSE
## kurtosis_roll_belt	1921.600000		2.02323922	FALSE	TRUE
## kurtosis_pitch_belt	600.500000		1.61553358	FALSE	TRUE
## kurtosis_yaw_belt	47.330049		0.01019264	FALSE	TRUE
## skewness_roll_belt	2135.111111		2.01304658	FALSE	TRUE
## skewness_roll_belt.1	600.500000		1.72255631	FALSE	TRUE
## skewness_yaw_belt	47.330049		0.01019264	FALSE	TRUE
## max_roll_belt		1.000000	0.99378249	FALSE	FALSE
## max_pitch_belt		1.538462	0.11211905	FALSE	FALSE
## max_yaw_belt		640.533333	0.34654979	FALSE	TRUE

```
training01 <- trainRaw[, !NZV$nzv]
testing01 <- testRaw[, !NZV$nzv]
dim(training01)
```

```
## [1] 19622 100
```

```
dim(testing01)
```

```
## [1] 20 100
```

2. Removing some columns of the dataset that do not contribute much to the accelerometer measurements.

```
regex <- grepl("^X|timestamp|user_name", names(training01))
training <- training01[, !regex]
testing <- testing01[, !regex]
dim(training)
```

```
## [1] 19622 95
```

```
dim(testing)
```

```
## [1] 20 95
```


3. Removing columns that contain NA's.

```
cond <- (colSums(is.na(training)) == 0)
training <- training[, cond]
testing <- testing[, cond]
```

Now, the cleaned training data set contains `{r} dim(training)[1]` observations and `{r} dim(training)[2]` variables, while the testing data set contains `{r} dim(testing)[1]` observations and `{r} dim(testing)[2]` variables.

Correlation Matrix of Columns in the Training Data set.

```
corrplot(cor(training[, -length(names(training))]), method = "color", tl.cex = 0.5)
```



Partitioning Training Set

we split the cleaned training set into a pure training data set (70%) and a validation data set (30%). We will use the validation data set to conduct cross validation in future steps.

```
set.seed(56789) # For reproducible purpose
inTrain <- createDataPartition(training$classe, p = 0.70, list = FALSE)
validation <- training[-inTrain, ]
training <- training[inTrain, ]
```

The Dataset now consists of `{r} dim(training)[2]` variables with the observations divided as following:

1. Training Data: `{r} dim(training)[1]` observations.

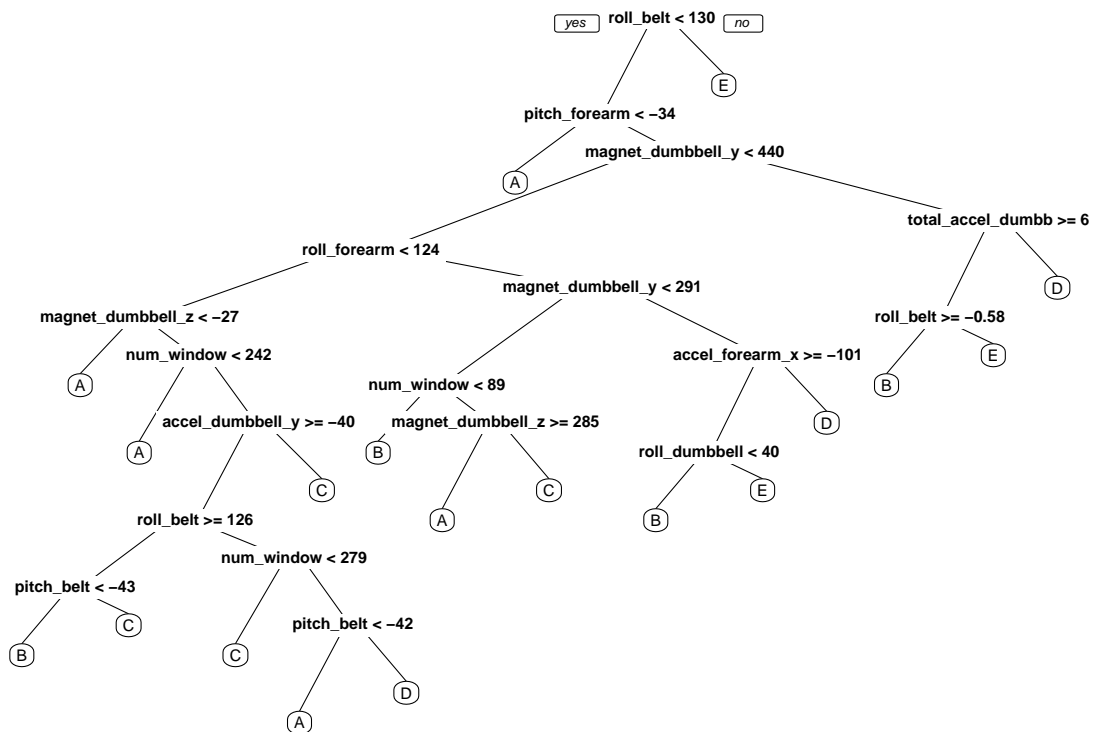
2. Validation Data: {r} dim(validation)[1] observations.
3. Testing Data: {r} dim(testing)[1] observations.

Data Modelling

Decision Tree

We fit a predictive model for activity recognition using Decision Tree algorithm.

```
modelTree <- rpart(classe ~ ., data = training, method = "class")
prp(modelTree)
```



Now, we estimate the performance of the model on the validation data set.

```
predictTree <- predict(modelTree, validation, type = "class")
confusionMatrix(validation$classe, predictTree)
```

Confusion Matrix and Statistics

##

Reference

Prediction A B C D E

A 1492 37 10 84 51

B 270 551 120 134 64

C 55 32 818 49 72

D 116 17 117 655 59

```
##           E    84    89    61   140   708
##
## Overall Statistics
##
##           Accuracy : 0.7178
##           95% CI : (0.7061, 0.7292)
##           No Information Rate : 0.3427
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.6409
##
## McNemar's Test P-Value : < 2.2e-16
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.7397  0.75895  0.7265  0.6168  0.7421
## Specificity      0.9529  0.88602  0.9563  0.9359  0.9242
## Pos Pred Value   0.8913  0.48376  0.7973  0.6795  0.6543
## Neg Pred Value   0.8753  0.96313  0.9366  0.9173  0.9488
## Prevalence       0.3427  0.12336  0.1913  0.1805  0.1621
## Detection Rate   0.2535  0.09363  0.1390  0.1113  0.1203
## Detection Prevalence 0.2845  0.19354  0.1743  0.1638  0.1839
## Balanced Accuracy 0.8463  0.82249  0.8414  0.7763  0.8331
```

```
accuracy <- postResample(predictTree, validation$classe)
ose <- 1 - as.numeric(confusionMatrix(validation$classe, predictTree)$overall[1])
rm(predictTree)
rm(modelTree)
```

The Estimated Accuracy of the Random Forest Model is “{r} accuracy[1]100% and the Estimated Out-of-Sample Error is {r} ose100“%.

Random Forest

We fit a predictive model for activity recognition using Random Forest algorithm because it automatically selects important variables and is robust to correlated covariates & outliers in general.

We will use 5-fold cross validation when applying the algorithm.

```
modelRF <- train(classe ~ ., data = training, method = "rf", trControl = trainControl(method = "cv", 5)
modelRF
```

```
## Random Forest
##
## 13737 samples
## 53 predictor
## 5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 10988, 10990, 10991, 10990, 10989
## Resampling results across tuning parameters:
```

```
##
## mtry Accuracy Kappa
## 2 0.9933033 0.9915283
## 27 0.9971614 0.9964095
## 53 0.9938853 0.9922657
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 27.
```

Now, we estimate the performance of the model on the validation data set.

```
predictRF <- predict(modelRF, validation)
confusionMatrix(validation$classe, predictRF)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    A    B    C    D    E
##           A 1674    0    0    0    0
##           B    1 1137    1    0    0
##           C    0    1 1025    0    0
##           D    0    0    0  964    0
##           E    0    0    0    2 1080
##
## Overall Statistics
##
##           Accuracy : 0.9992
##           95% CI : (0.998, 0.9997)
##           No Information Rate : 0.2846
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9989
##
## Mcnemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9994  0.9991  0.9990  0.9979  1.0000
## Specificity      1.0000  0.9996  0.9998  1.0000  0.9996
## Pos Pred Value   1.0000  0.9982  0.9990  1.0000  0.9982
## Neg Pred Value   0.9998  0.9998  0.9998  0.9996  1.0000
## Prevalence       0.2846  0.1934  0.1743  0.1641  0.1835
## Detection Rate   0.2845  0.1932  0.1742  0.1638  0.1835
## Detection Prevalence 0.2845  0.1935  0.1743  0.1638  0.1839
## Balanced Accuracy 0.9997  0.9993  0.9994  0.9990  0.9998
```

```
accuracy <- postResample(predictRF, validation$classe)
ose <- 1 - as.numeric(confusionMatrix(validation$classe, predictRF)$overall[1])
```

The Estimated Accuracy of the Random Forest Model is {r} accuracy[1]*100% and the Estimated Out-of-Sample Error is {r} ose*100%.

Random Forests yielded better Results, as expected!

Predicting The Manner of Exercise for Test Data Set

Now, we apply the Random Forest model to the original testing data set downloaded from the data source.

```
rm(accuracy)
rm(ose)
predict(modelRF, testing[, -length(names(testing))])
```

```
## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```

Generating Files to submit as answers for the Assignment

Function to generate files with predictions to submit for assignment.

```
pml_write_files = function(x){
  n = length(x)
  for(i in 1:n){
    filename = paste0("./Assignment_Solutions/problem_id_",i,".txt")
    write.table(x[i], file = filename, quote = FALSE, row.names = FALSE, col.names = FALSE)
  }
}
```

Generating the Files.

```
pml_write_files(predict(modelRF, testing[, -length(names(testing))]))
```